

# Hash tables, Binary search trees, Red-black trees

Congduan Li

Chinese University of Hong Kong, Shenzhen

*congduan.li@gmail.com*

Sep 26 & 28, 2017

# Direct-address tables

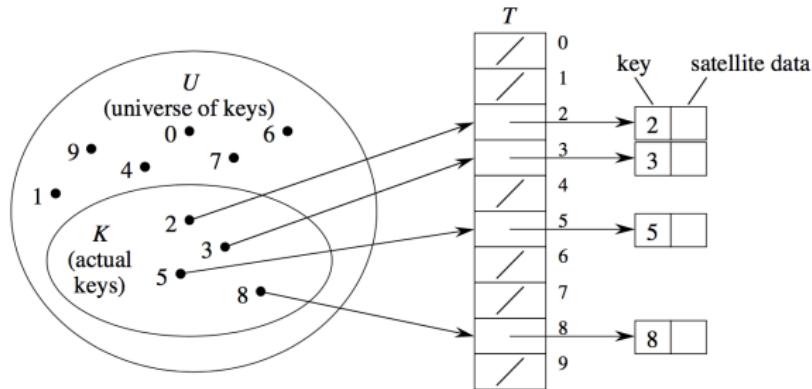
## **Scenario:**

- Maintain a dynamic set.
- Each element has a key drawn from a universe  $U = \{0, 1, \dots, m - 1\}$  where  $m$  isn't too large.
- No two elements have the same key.

Represent by a ***direct-address table***, or array,  $T[0 \dots m - 1]$ :

- Each ***slot***, or position, corresponds to a key in  $U$ .
- If there's an element  $x$  with key  $k$ , then  $T[k]$  contains a pointer to  $x$ .
- Otherwise,  $T[k]$  is empty, represented by NIL.

# Direct-address tables



Dictionary operations are trivial and take  $O(1)$  time each:

**DIRECT-ADDRESS-SEARCH( $T, k$ )**

**return  $T[k]$**

**DIRECT-ADDRESS-INSERT( $T, x$ )**

$T[key[x]] \leftarrow x$

**DIRECT-ADDRESS-DELETE( $T, x$ )**

$T[key[x]] \leftarrow \text{NIL}$

# Hash tables

The problem with direct addressing is if the universe  $U$  is large, storing a table of size  $|U|$  may be impractical or impossible.

Often, the set  $K$  of keys actually stored is small, compared to  $U$ , so that most of the space allocated for  $T$  is wasted.

- When  $K$  is much smaller than  $U$ , a hash table requires much less space than a direct-address table.
- Can reduce storage requirements to  $\Theta(|K|)$ .
- Can still get  $O(1)$  search time, but in the *average case*, not the *worst case*.

**Idea:** Instead of storing an element with key  $k$  in slot  $k$ , use a function  $h$  and store the element in slot  $h(k)$ .

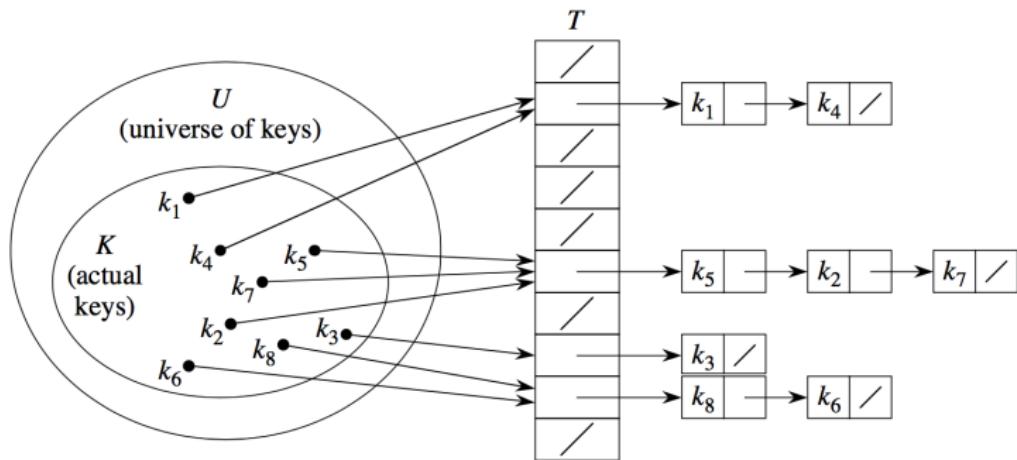
- We call  $h$  a **hash function**.
- $h : U \rightarrow \{0, 1, \dots, m - 1\}$ , so that  $h(k)$  is a legal slot number in  $T$ .
- We say that  $k$  **hashes** to slot  $h(k)$ .

# Collisions

***Collisions:*** When two or more keys hash to the same slot.

- Can happen when there are more possible keys than slots ( $|U| > m$ ).
- For a given set  $K$  of keys with  $|K| \leq m$ , may or may not happen. Definitely happens if  $|K| > m$ .
- Therefore, must be prepared to handle collisions in all cases.
- Use two methods: chaining and open addressing.
- Chaining is usually better than open addressing.

# Collision resolution by chaining



[This figure shows singly linked lists. If we want to delete elements, it's better to use doubly linked lists.]

- Slot  $j$  contains a pointer to the head of the list of all stored elements that hash to  $j$  [or to the sentinel if using a circular, doubly linked list with a sentinel] ,
- If there are no such elements, slot  $j$  contains NIL.

# Operations with chaining

- ***Insertion:***

CHAINED-HASH-INSERT( $T, x$ )

insert  $x$  at the head of list  $T[h(\text{key}[x])]$

- Worst-case running time is  $O(1)$ .
- Assumes that the element being inserted isn't already in the list.
- It would take an additional search to check if it was already inserted.

- ***Search:***

CHAINED-HASH-SEARCH( $T, k$ )

search for an element with key  $k$  in list  $T[h(k)]$

Running time is proportional to the length of the list of elements in slot  $h(k)$ .

# Operations with chaining

- ***Deletion:***

CHAINED-HASH-DELETE( $T, x$ )

delete  $x$  from the list  $T[h(\text{key}[x])]$

- Given pointer  $x$  to the element to delete, so no search is needed to find this element.
- Worst-case running time is  $O(1)$  time if the lists are doubly linked.
- If the lists are singly linked, then deletion takes as long as searching, because we must find  $x$ 's predecessor in its list in order to correctly update *next* pointers.

# Analysis of hashing with chaining

Given a key, how long does it take to find an element with that key, or to determine that there is no element with that key?

- Analysis is in terms of the ***load factor***  $\alpha = n/m$ :
  - $n$  = # of elements in the table.
  - $m$  = # of slots in the table = # of (possibly empty) linked lists.
  - Load factor is average number of elements per linked list.
  - Can have  $\alpha < 1$ ,  $\alpha = 1$ , or  $\alpha > 1$ .
- Worst case is when all  $n$  keys hash to the same slot  $\Rightarrow$  get a single list of length  $n$   
 $\Rightarrow$  worst-case time to search is  $\Theta(n)$ , plus time to compute hash function.
- Average case depends on how well the hash function distributes the keys among the slots.

# Analysis of hashing with chaining

We focus on average-case performance of hashing with chaining.

- Assume ***simple uniform hashing***: any given element is equally likely to hash into any of the  $m$  slots.
- For  $j = 0, 1, \dots, m - 1$ , denote the length of list  $T[j]$  by  $n_j$ . Then  $n = n_0 + n_1 + \dots + n_{m-1}$ .
- Average value of  $n_j$  is  $E[n_j] = \alpha = n/m$ .
- Assume that we can compute the hash function in  $O(1)$  time, so that the time required to search for the element with key  $k$  depends on the length  $n_{h(k)}$  of the list  $T[h(k)]$ .

We consider two cases:

- If the hash table contains no element with key  $k$ , then the search is unsuccessful.
- If the hash table does contain an element with key  $k$ , then the search is successful.

# Unsuccessful search

## **Theorem**

An unsuccessful search takes expected time  $\Theta(1 + \alpha)$ .

**Proof** Simple uniform hashing  $\Rightarrow$  any key not already in the table is equally likely to hash to any of the  $m$  slots.

To search unsuccessfully for any key  $k$ , need to search to the end of the list  $T[h(k)]$ . This list has expected length  $E[n_{h(k)}] = \alpha$ . Therefore, the expected number of elements examined in an unsuccessful search is  $\alpha$ .

Adding in the time to compute the hash function, the total time required is  $\Theta(1 + \alpha)$ . ■

# Successful search

## Theorem

A successful search takes expected time  $\Theta(1 + \alpha)$ .

**Proof** Assume that the element  $x$  being searched for is equally likely to be any of the  $n$  elements stored in the table.

The number of elements examined during a successful search for  $x$  is 1 more than the number of elements that appear before  $x$  in  $x$ 's list. These are the elements inserted *after*  $x$  was inserted (because we insert at the head of the list).

So we need to find the average, over the  $n$  elements  $x$  in the table, of how many elements were inserted into  $x$ 's list after  $x$  was inserted.

For  $i = 1, 2, \dots, n$ , let  $x_i$  be the  $i$ th element inserted into the table, and let  $k_i = \text{key}[x_i]$ .

For all  $i$  and  $j$ , define indicator random variable  $X_{ij} = I\{h(k_i) = h(k_j)\}$ .

Simple uniform hashing  $\Rightarrow \Pr\{h(k_i) = h(k_j)\} = 1/m \Rightarrow E[X_{ij}] = 1/m$

# Successful search

Expected number of elements examined in a successful search is

$$\begin{aligned} & E \left[ \frac{1}{n} \sum_{i=1}^n \left( 1 + \sum_{j=i+1}^n X_{ij} \right) \right] \\ &= \frac{1}{n} \sum_{i=1}^n \left( 1 + \sum_{j=i+1}^n E[X_{ij}] \right) \quad (\text{linearity of expectation}) \\ &= \frac{1}{n} \sum_{i=1}^n \left( 1 + \sum_{j=i+1}^n \frac{1}{m} \right) \\ &= 1 + \frac{1}{nm} \sum_{i=1}^n (n - i) \\ &= 1 + \frac{1}{nm} \left( \sum_{i=1}^n n - \sum_{i=1}^n i \right) \\ &= 1 + \frac{1}{nm} \left( n^2 - \frac{n(n+1)}{2} \right) \quad (\text{equation (A.1)}) \\ &= 1 + \frac{n-1}{2m} \\ &= 1 + \frac{\alpha}{2} - \frac{\alpha}{2n}. \end{aligned}$$

Adding in the time for computing the hash function, we get that the expected total time for a successful search is  $\Theta(2 + \alpha/2 - \alpha/2n) = \Theta(1 + \alpha)$ .

# Successful search

**Interpretation:** If  $n = O(m)$ , then  $\alpha = n/m = O(m)/m = O(1)$ , which means that searching takes constant time on average.

Since insertion takes  $O(1)$  worst-case time and deletion takes  $O(1)$  worst-case time when the lists are doubly linked, all dictionary operations take  $O(1)$  time on average.

# Hash functions

## What makes a good hash function?

- Ideally, the hash function satisfies the assumption of simple uniform hashing.
- In practice, it's not possible to satisfy this assumption, since we don't know in advance the probability distribution that keys are drawn from, and the keys may not be drawn independently.
- Often use heuristics, based on the domain of the keys, to create a hash function that performs well.

## Keys as natural numbers

- Hash functions assume that the keys are natural numbers.
- When they're not, have to interpret them as natural numbers.
- **Example:** Interpret a character string as an integer expressed in some radix notation. Suppose the string is CLRS:
  - ASCII values: C = 67, L = 76, R = 82, S = 83.
  - There are 128 basic ASCII values.
  - So interpret CLRS as  $(67 \cdot 128^3) + (76 \cdot 128^2) + (82 \cdot 128^1) + (83 \cdot 128^0) = 141,764,947$ .

# Hash functions

## Division method

$$h(k) = k \bmod m .$$

**Example:**  $m = 20$  and  $k = 91 \Rightarrow h(k) = 11$ .

**Advantage:** Fast, since requires just one division operation.

**Disadvantage:** Have to avoid certain values of  $m$ :

- Powers of 2 are bad. If  $m = 2^p$  for integer  $p$ , then  $h(k)$  is just the least significant  $p$  bits of  $k$ .
- If  $k$  is a character string interpreted in radix  $2^p$  (as in CLRS example), then  $m = 2^p - 1$  is bad: permuting characters in a string does not change its hash value (Exercise 11.3-3).

**Good choice for  $m$ :** A prime not too close to an exact power of 2.

# Hash functions

## Multiplication method

1. Choose constant  $A$  in the range  $0 < A < 1$ .
2. Multiply key  $k$  by  $A$ .
3. Extract the fractional part of  $kA$ .
4. Multiply the fractional part by  $m$ .
5. Take the floor of the result.

Put another way,  $h(k) = \lfloor m(kA \bmod 1) \rfloor$ , where  $kA \bmod 1 = kA - \lfloor kA \rfloor =$  fractional part of  $kA$ .

**Disadvantage:** Slower than division method.

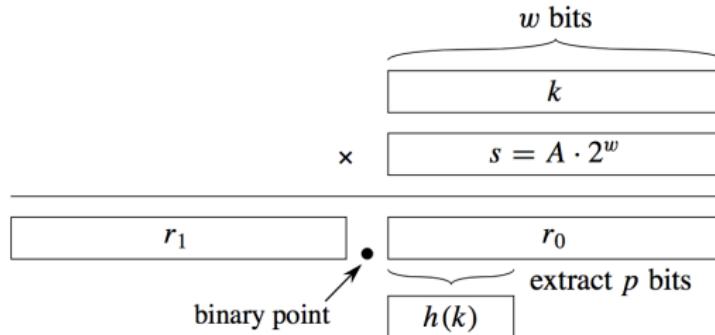
**Advantage:** Value of  $m$  is not critical.

**(Relatively) easy implementation:**

- Choose  $m = 2^p$  for some integer  $p$ .
- Let the word size of the machine be  $w$  bits.
- Assume that  $k$  fits into a single word. ( $k$  takes  $w$  bits.)
- Let  $s$  be an integer in the range  $0 < s < 2^w$ . ( $s$  takes  $w$  bits.)

# Hash functions

- Restrict  $A$  to be of the form  $s/2^w$ .



- Multiply  $k$  by  $s$ .
- Since we're multiplying two  $w$ -bit words, the result is  $2w$  bits,  $r_1 2^w + r_0$ , where  $r_1$  is the high-order word of the product and  $r_0$  is the low-order word.
- $r_1$  holds the integer part of  $kA$  ( $\lfloor kA \rfloor$ ) and  $r_0$  holds the fractional part of  $kA$  ( $kA \bmod 1 = kA - \lfloor kA \rfloor$ ). Think of the “binary point” (analog of decimal point, but for binary representation) as being between  $r_1$  and  $r_0$ . Since we don't care about the integer part of  $kA$ , we can forget about  $r_1$  and just use  $r_0$ .
- Since we want  $\lfloor m(kA \bmod 1) \rfloor$ , we could get that value by shifting  $r_0$  to the left by  $p = \lg m$  bits and then taking the  $p$  bits that were shifted to the left of the binary point.

# Hash functions

- We don't need to shift. The  $p$  bits that would have been shifted to the left of the binary point are the  $p$  most significant bits of  $r_0$ . So we can just take these bits after having formed  $r_0$  by multiplying  $k$  by  $s$ .
- **Example:**  $m = 8$  (implies  $p = 3$ ),  $w = 5$ ,  $k = 21$ . Must have  $0 < s < 2^5$ ; choose  $s = 13 \Rightarrow A = 13/32$ .
  - Using just the formula to compute  $h(k)$ :  $kA = 21 \cdot 13/32 = 273/32 = 8\frac{17}{32}$   
 $\Rightarrow kA \bmod 1 = 17/32 \Rightarrow m(kA \bmod 1) = 8 \cdot 17/32 = 17/4 = 4\frac{1}{4} \Rightarrow \lfloor m(kA \bmod 1) \rfloor = 4$ , so that  $h(k) = 4$ .
  - Using the implementation:  $ks = 21 \cdot 13 = 273 = 8 \cdot 2^5 + 17 \Rightarrow r_1 = 8$ ,  $r_0 = 17$ . Written in  $w = 5$  bits,  $r_0 = 10001$ . Take the  $p = 3$  most significant bits of  $r_0$ , get 100 in binary, or 4 in decimal, so that  $h(k) = 4$ .

## How to choose A:

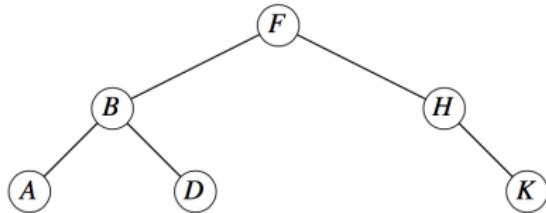
- The multiplication method works with any legal value of  $A$ .
- But it works better with some values than with others, depending on the keys being hashed.
- Knuth suggests using  $A \approx (\sqrt{5} - 1)/2$ .

# Binary search trees (Chap 12)

Binary search trees are an important data structure for dynamic sets.

- Accomplish many dynamic-set operations in  $O(h)$  time, where  $h$  = height of tree.
- As in Section 10.4, we represent a binary tree by a linked data structure in which each node is an object.
- $\text{root}[T]$  points to the root of tree  $T$ .
- Each node contains the fields
  - $\text{key}$  (and possibly other satellite data).
  - $\text{left}$ : points to left child.
  - $\text{right}$ : points to right child.
  - $p$ : points to parent.  $p[\text{root}[T]] = \text{NIL}$ .
- Stored keys must satisfy the ***binary-search-tree property***.
  - If  $y$  is in left subtree of  $x$ , then  $\text{key}[y] \leq \text{key}[x]$ .
  - If  $y$  is in right subtree of  $x$ , then  $\text{key}[y] \geq \text{key}[x]$ .

# Binary search trees



The binary-search-tree property allows us to print keys in a binary search tree in order, recursively, using an algorithm called an *inorder tree walk*. Elements are printed in monotonically increasing order.

How INORDER-TREE-WALK works:

- Check to make sure that  $x$  is not NIL.
- Recursively, print the keys of the nodes in  $x$ 's left subtree.
- Print  $x$ 's key.
- Recursively, print the keys of the nodes in  $x$ 's right subtree.

# Binary search trees

INORDER-TREE-WALK( $x$ )

**if**  $x \neq \text{NIL}$

**then** INORDER-TREE-WALK( $\text{left}[x]$ )

print  $\text{key}[x]$

INORDER-TREE-WALK( $\text{right}[x]$ )

**Example:** Do the inorder tree walk on the example above, getting the output  $ABDFHK$ .

**Correctness:** Follows by induction directly from the binary-search-tree property.

**Time:** Intuitively, the walk takes  $\Theta(n)$  time for a tree with  $n$  nodes, because we visit and print each node once. [Book has formal proof.]

# Search on binary search trees

```
TREE-SEARCH( $x, k$ )
if  $x = \text{NIL}$  or  $k = \text{key}[x]$ 
    then return  $x$ 
if  $k < \text{key}[x]$ 
    then return TREE-SEARCH( $\text{left}[x], k$ )
else return TREE-SEARCH( $\text{right}[x], k$ )
```

Initial call is TREE-SEARCH( $\text{root}[T], k$ ).

**Example:** Search for values  $D$  and  $C$  in the example tree from above.

**Time:** The algorithm recurses, visiting nodes on a downward path from the root. Thus, running time is  $O(h)$ , where  $h$  is the height of the tree.

# Max and Min on binary search trees

The binary-search-tree property guarantees that

- the minimum key of a binary search tree is located at the leftmost node, and
- the maximum key of a binary search tree is located at the rightmost node.

Traverse the appropriate pointers (*left* or *right*) until NIL is reached.

**TREE-MINIMUM( $x$ )**

```
while left[ $x$ ] ≠ NIL
    do  $x \leftarrow \text{left}[x]$ 
return  $x$ 
```

**TREE-MAXIMUM( $x$ )**

```
while right[ $x$ ] ≠ NIL
    do  $x \leftarrow \text{right}[x]$ 
return  $x$ 
```

**Time:** Both procedures visit nodes that form a downward path from the root to a leaf. Both procedures run in  $O(h)$  time, where  $h$  is the height of the tree.

# Successor and predecessor

There are two cases:

1. If node  $x$  has a non-empty right subtree, then  $x$ 's successor is the minimum in  $x$ 's right subtree.
2. If node  $x$  has an empty right subtree, notice that:
  - As long as we move to the left up the tree (move up through right children), we're visiting smaller keys.
  - $x$ 's successor  $y$  is the node that  $x$  is the predecessor of ( $x$  is the maximum in  $y$ 's left subtree).

TREE-SUCCESSOR( $x$ )

**if**  $right[x] \neq \text{NIL}$   
**then return** TREE-MINIMUM( $right[x]$ )

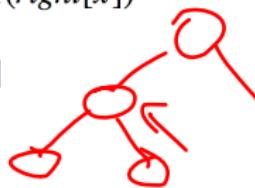
$y \leftarrow p[x]$

**while**  $y \neq \text{NIL}$  and  $x = right[y]$

**do**  $x \leftarrow y$

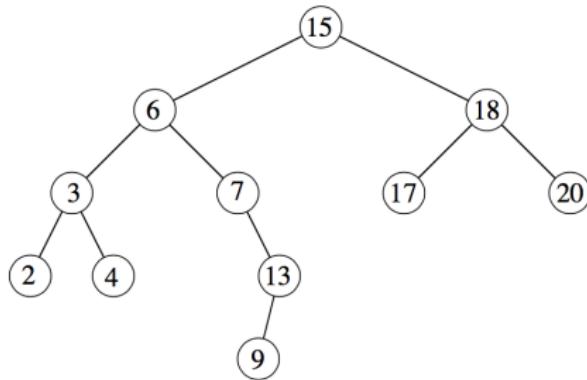
$y \leftarrow p[y]$

**return**  $y$  (Y=nil, 没有)



TREE-PREDECESSOR is symmetric to TREE-SUCCESSOR.

# Successor and predecessor



- Find the successor of the node with key value 15. (Answer: Key value 17)
- Find the successor of the node with key value 6. (Answer: Key value 7)
- Find the successor of the node with key value 4. (Answer: Key value 6)
- Find the predecessor of the node with key value 6. (Answer: Key value 4)

**Time:** For both the TREE-SUCCESSOR and TREE-PREDECESSOR procedures, in both cases, we visit nodes on a path down the tree or up the tree. Thus, running time is  $O(h)$ , where  $h$  is the height of the tree.

# Insertion

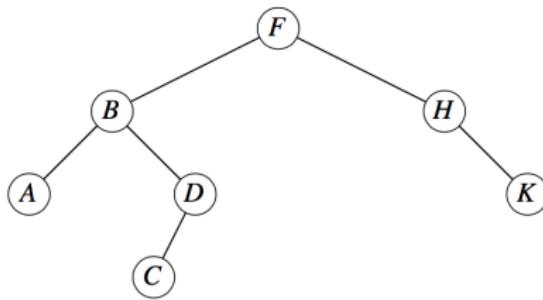
```
TREE-INSERT( $T, z$ )
 $y \leftarrow \text{NIL}$ 
 $x \leftarrow \text{root}[T]$ 
while  $x \neq \text{NIL}$ 
    do  $y \leftarrow x$ 
        if  $\text{key}[z] < \text{key}[x]$ 
            then  $x \leftarrow \text{left}[x]$ 
            else  $x \leftarrow \text{right}[x]$ 
 $p[z] \leftarrow y$ 
if  $y = \text{NIL}$ 
    then  $\text{root}[T] \leftarrow z$             $\triangleright$  Tree  $T$  was empty
    else if  $\text{key}[z] < \text{key}[y]$ 
        then  $\text{left}[y] \leftarrow z$ 
        else  $\text{right}[y] \leftarrow z$ 
```

- To insert value  $v$  into the binary search tree, the procedure is given node  $z$ , with  $\text{key}[z] = v$ ,  $\text{left}[z] = \text{NIL}$ , and  $\text{right}[z] = \text{NIL}$ .
- Beginning at root of the tree, trace a downward path, maintaining two pointers.
  - Pointer  $x$ : traces the downward path.
  - Pointer  $y$ : “trailing pointer” to keep track of parent of  $x$ .

# Insertion

- Traverse the tree downward by comparing the value of node at  $x$  with  $v$ , and move to the left or right child accordingly.
- When  $x$  is NIL, it is at the correct position for node  $z$ .
- Compare  $z$ 's value with  $y$ 's value, and insert  $z$  at either  $y$ 's *left* or *right*, appropriately.

**Example:** Run TREE-INSERT( $C$ ) on the first sample binary search tree. Result:



**Time:** Same as TREE-SEARCH. On a tree of height  $h$ , procedure takes  $O(h)$  time.

# Deletion

TREE-DELETE is broken into three cases.

**Case 1:**  $z$  has no children.

- Delete  $z$  by making the parent of  $z$  point to NIL, instead of to  $z$ .

**Case 2:**  $z$  has one child.

- Delete  $z$  by making the parent of  $z$  point to  $z$ 's child, instead of to  $z$ .

**Case 3:**  $z$  has two children.

- $z$ 's successor  $y$  has either no children or one child. ( $y$  is the minimum node—with no left child—in  $z$ 's right subtree.)
- Delete  $y$  from the tree (via Case 1 or 2).
- Replace  $z$ 's key and satellite data with  $y$ 's.

# Deletion

TREE-DELETE( $T, z$ )

▷ Determine which node  $y$  to splice out: either  $z$  or  $z$ 's successor.

**if**  $left[z] = NIL$  **or**  $right[z] = NIL$

**then**  $y \leftarrow z$

**else**  $y \leftarrow \text{TREE-SUCCESSOR}(z)$

▷  $x$  is set to a non-NIL child of  $y$ , or to NIL if  $y$  has no children.

**if**  $left[y] \neq NIL$

**then**  $x \leftarrow left[y]$

**else**  $x \leftarrow right[y]$

▷  $y$  is removed from the tree by manipulating pointers of  $p[y]$  and  $x$ .

**if**  $x \neq NIL$

**then**  $p[x] \leftarrow p[y]$

**if**  $p[y] = NIL$

**then**  $root[T] \leftarrow x$

**else if**  $y = left[p[y]]$

**then**  $left[p[y]] \leftarrow x$

**else**  $right[p[y]] \leftarrow x$

▷ If it was  $z$ 's successor that was spliced out, copy its data into  $z$ .

**if**  $y \neq z$

**then**  $key[z] \leftarrow key[y]$

        copy  $y$ 's satellite data into  $z$

**return**  $y$

# Deletion

**Example:** We can demonstrate on the above sample tree.

- For Case 1, delete  $K$ .
- For Case 2, delete  $H$ .
- For Case 3, delete  $B$ , swapping it with  $C$ .

**Time:**  $O(h)$ , on a tree of height  $h$ .

We've been analyzing running time in terms of  $h$  (the height of the binary search tree), instead of  $n$  (the number of nodes in the tree).

- Problem: Worst case for binary search tree is  $\Theta(n)$ —no better than linked list.
- Solution: Guarantee small height (balanced tree)— $h = O(\lg n)$ .

In later chapters, by varying the properties of binary search trees, we will be able to analyze running time in terms of  $n$ .

- Method: Restructure the tree if necessary. Nothing special is required for querying, but there may be extra work when changing the structure of the tree (inserting or deleting).

Red-black trees are a special class of binary trees that avoids the worst-case behavior of  $O(n)$  like “plain” binary search trees. Red-black trees are covered in detail in Chapter 13.

# Red-black trees (Chap 13)

A **red-black tree** is a binary search tree + 1 bit per node: an attribute *color*, which is either red or black.

All leaves are empty (*nil*) and colored black.

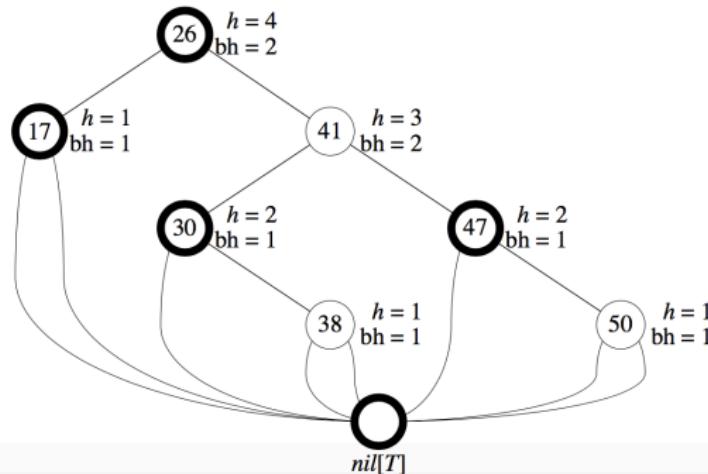
- We use a single sentinel,  $\text{nil}[T]$ , for all the leaves of red-black tree  $T$ .
- $\text{color}[\text{nil}[T]]$  is black.
- The root's parent is also  $\text{nil}[T]$ .

All other attributes of binary search trees are inherited by red-black trees (*key*, *left*, *right*, and *p*). We don't care about the key in  $\text{nil}[T]$ .

# Red-black trees properties

1. Every node is either red or black.
2. The root is black.
3. Every leaf ( $\text{nil}[T]$ ) is black.
4. If a node is red, then both its children are black. (Hence no two reds in a row on a simple path from the root to a leaf.)
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes.

Example:



# Height of red-black trees

- **Height of a node** is the number of edges in a longest path to a leaf.
- **Black-height** of a node  $x$ :  $\text{bh}(x)$  is the number of black nodes (including  $\text{nil}[T]$ ) on the path from  $x$  to leaf, not counting  $x$ . By property 5, black-height is well defined.

[Now label the example tree with height  $h$  and bh values.]

### **Claim**

Any node with height  $h$  has black-height  $\geq h/2$ .

**Proof** By property 4,  $\leq h/2$  nodes on the path from the node to a leaf are red. Hence  $\geq h/2$  are black. ■ (claim)

# Height of red-black trees

## Claim

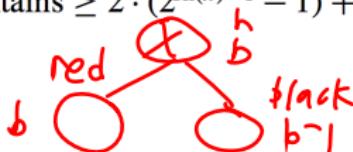
The subtree rooted at any node  $x$  contains  $\geq 2^{\text{bh}(x)} - 1$  internal nodes.

**Proof** By induction on height of  $x$ .

**Basis:** Height of  $x = 0 \Rightarrow x$  is a leaf  $\Rightarrow \text{bh}(x) = 0$ . The subtree rooted at  $x$  has 0 internal nodes.  $2^0 - 1 = 0$ .

**Inductive step:** Let the height of  $x$  be  $h$  and  $\text{bh}(x) = b$ . Any child of  $x$  has height  $h - 1$  and black-height either  $b$  (if the child is red) or  $b - 1$  (if the child is black). By the inductive hypothesis, each child has  $\geq 2^{\text{bh}(x)-1} - 1$  internal nodes. Thus, the subtree rooted at  $x$  contains  $\geq 2 \cdot (2^{\text{bh}(x)-1} - 1) + 1 = 2^{\text{bh}(x)} - 1$  internal nodes. (The +1 is for  $x$  itself.) ■ (claim)

## Lemma



A red-black tree with  $n$  internal nodes has height  $\leq 2 \lg(n + 1)$ .

**Proof** Let  $h$  and  $b$  be the height and black-height of the root, respectively. By the above two claims,

$$n \geq 2^b - 1 \geq 2^{h/2} - 1. \quad (b \geq 2^{h/2})$$

Adding 1 to both sides and then taking logs gives  $\lg(n + 1) \geq h/2$ , which implies that  $h \leq 2 \lg(n + 1)$ . ■ (theorem)

# Operations on red-black trees

The non-modifying binary-search-tree operations MINIMUM, MAXIMUM, SUCCESSOR, PREDECESSOR, and SEARCH run in  $O(\text{height})$  time. Thus, they take  $O(\lg n)$  time on red-black trees.

Insertion and deletion are not so easy.

If we insert, what color to make the new node?

*Color the new node in red and fix 4*

- Red? Might violate property 4 */2 if empty*
- Black? Might violate property 5.

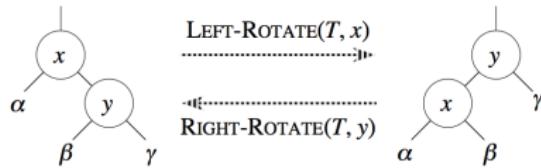
If we delete, thus removing a node, what color was the node that was removed?

- Red? OK, since we won't have changed any black-heights, nor will we have created two red nodes in a row. Also, cannot cause a violation of property 2, since if the removed node was red, it could not have been the root.
- Black? Could cause there to be two reds in a row (violating property 4), and can also cause a violation of property 5. Could also cause a violation of property 2, if the removed node was the root and its child—which becomes the new root—was red.

# Rotation

- The basic tree-restructuring operation.
- Needed to maintain red-black trees as balanced binary search trees.
- Changes the local pointer structure. (Only pointers are changed.)
- Won't upset the binary-search-tree property.
- Have both left rotation and right rotation. They are inverses of each other.
- A rotation takes a red-black-tree and a node within the tree.

# Left rotation



LEFT-ROTATE( $T, x$ )

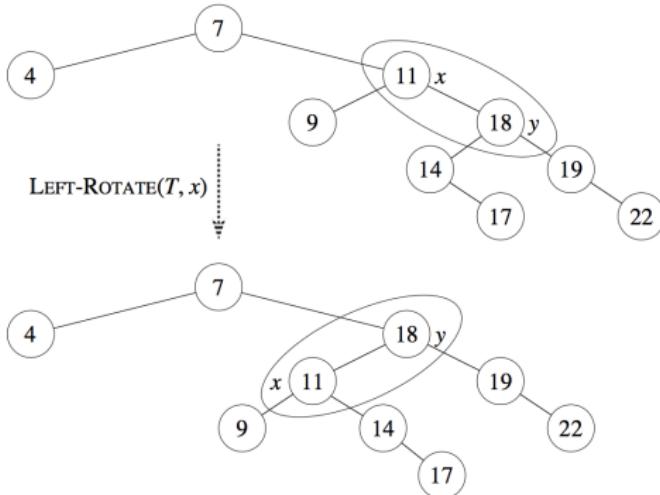
```
 $y \leftarrow \text{right}[x]$             $\triangleright$  Set  $y$ .  
 $\text{right}[x] \leftarrow \text{left}[y]$         $\triangleright$  Turn  $y$ 's left subtree into  $x$ 's right subtree.  
if  $\text{left}[y] \neq \text{nil}[T]$   
    then  $p[\text{left}[y]] \leftarrow x$   
 $p[y] \leftarrow p[x]$                    $\triangleright$  Link  $x$ 's parent to  $y$ .  
if  $p[x] = \text{nil}[T]$   
    then  $\text{root}[T] \leftarrow y$   
    else if  $x = \text{left}[p[x]]$   
        then  $\text{left}[p[x]] \leftarrow y$   
        else  $\text{right}[p[x]] \leftarrow y$   
 $\text{left}[y] \leftarrow x$                   $\triangleright$  Put  $x$  on  $y$ 's left.  
 $p[x] \leftarrow y$ 
```

The pseudocode for LEFT-ROTATE assumes that

- $\text{right}[x] \neq \text{nil}[T]$ , and
- root's parent is  $\text{nil}[T]$ .

Pseudocode for RIGHT-ROTATE is symmetric: exchange *left* and *right* everywhere.

# Example



- Before rotation: keys of  $x$ 's left subtree  $\leq 11 \leq$  keys of  $y$ 's left subtree  $\leq 18 \leq$  keys of  $y$ 's right subtree.
- Rotation makes  $y$ 's left subtree into  $x$ 's right subtree.
- After rotation: keys of  $x$ 's left subtree  $\leq 11 \leq$  keys of  $x$ 's right subtree  $\leq 18 \leq$  keys of  $y$ 's right subtree.

**Time:**  $O(1)$  for both LEFT-ROTATE and RIGHT-ROTATE, since a constant number of pointers are modified.

# RB insertion

```
RB-INSERT( $T, z$ )
   $y \leftarrow nil[T]$ 
   $x \leftarrow root[T]$ 
  while  $x \neq nil[T]$ 
    do  $y \leftarrow x$ 
      if  $key[z] < key[x]$ 
        then  $x \leftarrow left[x]$ 
        else  $x \leftarrow right[x]$ 
   $p[z] \leftarrow y$ 
  if  $y = nil[T]$ 
    then  $root[T] \leftarrow z$ 
    else if  $key[z] < key[y]$ 
      then  $left[y] \leftarrow z$ 
      else  $right[y] \leftarrow z$ 
   $left[z] \leftarrow nil[T]$ 
   $right[z] \leftarrow nil[T]$ 
   $color[z] \leftarrow RED$ 
  RB-INSERT-FIXUP( $T, z$ )
```

- RB-INSERT ends by coloring the new node  $z$  red.
- Then it calls RB-INSERT-FIXUP because we could have violated a red-black property.

# RB insertion fixup

Which property might be violated?

1. OK.
2. If  $z$  is the root, then there's a violation. Otherwise, OK.
3. OK.
4. If  $p[z]$  is red, there's a violation: both  $z$  and  $p[z]$  are red.
5. OK.

Remove the violation by calling RB-INSERT-FIXUP:

RB-INSERT-FIXUP( $T, z$ )

```
while color[p[z]] = RED
    do if p[z] = left[p[p[z]]]
        then y ← right[p[p[z]]]
            if color[y] = RED
                then color[p[z]] ← BLACK          ▷ Case 1
                    color[y] ← BLACK             ▷ Case 1
                    color[p[p[z]]] ← RED         ▷ Case 1
                    z ← p[p[z]]                 ▷ Case 1
                else if z = right[p[z]]
                    then z ← p[z]                  ▷ Case 2
                        LEFT-ROTATE(T, z)       ▷ Case 2
                    color[p[z]] ← BLACK          ▷ Case 3
                    color[p[p[z]]] ← RED         ▷ Case 3
                    RIGHT-ROTATE(T, p[p[z]])   ▷ Case 3
                else (same as then clause
                      with "right" and "left" exchanged)
                    color[root[T]] ← BLACK
```

# Correctness

## Loop invariant:

At the start of each iteration of the **while** loop,

- a.  $z$  is red.
- b. There is at most one red-black violation:
  - Property 2:  $z$  is a red root, or
  - Property 4:  $z$  and  $p[z]$  are both red.

*[The book has a third part of the loop invariant, but we omit it for lecture.]*

**Initialization:** We've already seen why the loop invariant holds initially.

**Termination:** The loop terminates because  $p[z]$  is black. Hence, property 4 is OK. Only property 2 might be violated, and the last line fixes it.

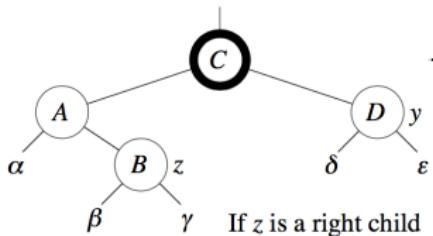
**Maintenance:** We drop out when  $z$  is the root (since then  $p[z]$  is the sentinel  $nil[T]$ , which is black). When we start the loop body, the only violation is of property 4.

There are 6 cases, 3 of which are symmetric to the other 3. The cases are not mutually exclusive. We'll consider cases in which  $p[z]$  is a left child.

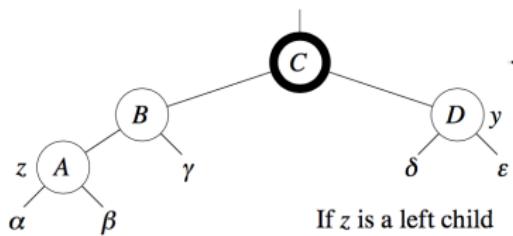
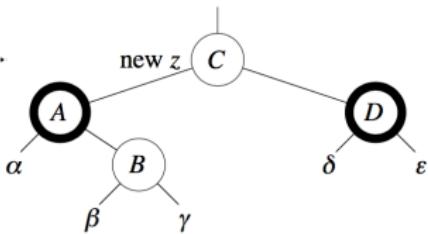
Let  $y$  be  $z$ 's uncle ( $p[z]$ 's sibling).

# Case 1

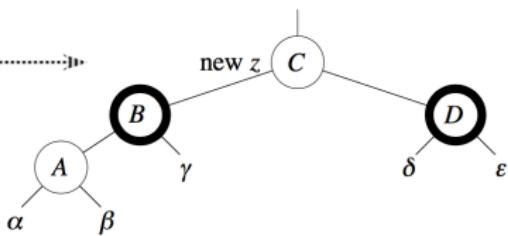
**Case 1:**  $y$  is red



If  $z$  is a right child



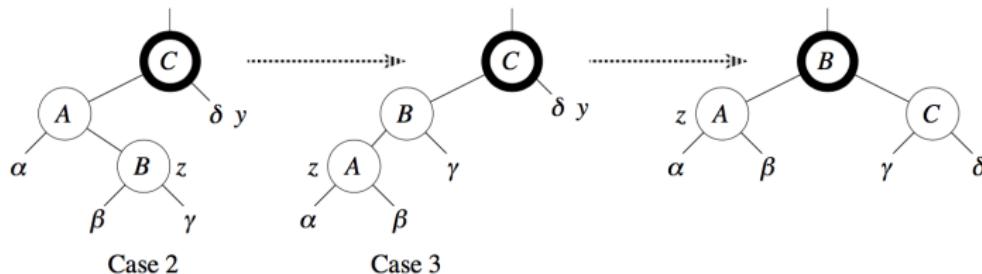
If  $z$  is a left child



- $p[p[z]]$  ( $z$ 's grandparent) must be black, since  $z$  and  $p[z]$  are both red and there are no other violations of property 4.
- Make  $p[z]$  and  $y$  black  $\Rightarrow$  now  $z$  and  $p[z]$  are not both red. But property 5 might now be violated.
- Make  $p[p[z]]$  red  $\Rightarrow$  restores property 5.
- The next iteration has  $p[p[z]]$  as the new  $z$  (i.e.,  $z$  moves up 2 levels).

## Case 2 and Case 3

**Case 2:**  $y$  is black,  $z$  is a right child



- Left rotate around  $p[z] \Rightarrow$  now  $z$  is a left child, and both  $z$  and  $p[z]$  are red.
- Takes us immediately to case 3.

**Case 3:**  $y$  is black,  $z$  is a left child

- Make  $p[z]$  black and  $p[p[z]]$  red.
- Then right rotate on  $p[p[z]]$ .
- No longer have 2 reds in a row.
- $p[z]$  is now black  $\Rightarrow$  no more iterations.

# Time analysis

## Analysis

$O(\lg n)$  time to get through RB-INSERT up to the call of RB-INSERT-FIXUP.

Within RB-INSERT-FIXUP:

- Each iteration takes  $O(1)$  time.
- Each iteration is either the last one or it moves  $z$  up 2 levels.
- $O(\lg n)$  levels  $\Rightarrow O(\lg n)$  time.
- Also note that there are at most 2 rotations overall.

Thus, insertion into a red-black tree takes  $O(\lg n)$  time.

# RB deletion

```
RB-DELETE( $T, z$ )
if  $left[z] = nil[T]$  or  $right[z] = nil[T]$ 
    then  $y \leftarrow z$ 
    else  $y \leftarrow \text{TREE-SUCCESSOR}(z)$ 
if  $left[y] \neq nil[T]$ 
    then  $x \leftarrow left[y]$ 
    else  $x \leftarrow right[y]$ 
 $p[x] \leftarrow p[y]$ 
if  $p[y] = nil[T]$ 
    then  $root[T] \leftarrow x$ 
    else if  $y = left[p[y]]$ 
        then  $left[p[y]] \leftarrow x$ 
        else  $right[p[y]] \leftarrow x$ 
if  $y \neq z$ 
    then  $key[z] \leftarrow key[y]$ 
        copy  $y$ 's satellite data into  $z$ 
if  $color[y] = \text{BLACK}$ 
    then RB-DELETE-FIXUP( $T, x$ )
return  $y$ 
```

- $y$  is the node that was actually spliced out.
- $x$  is either
  - $y$ 's sole non-sentinel child before  $y$  was spliced out, or
  - the sentinel, if  $y$  had no children.

# RB deletion

If  $y$  is black, we could have violations of red-black properties:

1. OK.
2. If  $y$  is the root and  $x$  is red, then the root has become red.
3. OK.
4. Violation if  $p[y]$  and  $x$  are both red.
5. Any path containing  $y$  now has 1 fewer black node.
  - Correct by giving  $x$  an “extra black.”
  - Add 1 to count of black nodes on paths containing  $x$ .
  - Now property 5 is OK, but property 1 is not.
  - $x$  is either ***doubly black*** (if  $\text{color}[x] = \text{BLACK}$ ) or ***red & black*** (if  $\text{color}[x] = \text{RED}$ ).
  - The attribute  $\text{color}[x]$  is still either RED or BLACK. No new values for  $\text{color}$  attribute.
  - In other words, the extra blackness on a node is by virtue of  $x$  pointing to the node.

# RB deletion fixup

RB-DELETE-FIXUP( $T, x$ )

**while**  $x \neq \text{root}[T]$  and  $\text{color}[x] = \text{BLACK}$

**do if**  $x = \text{left}[p[x]]$

**then**  $w \leftarrow \text{right}[p[x]]$

**if**  $\text{color}[w] = \text{RED}$

**then**  $\text{color}[w] \leftarrow \text{BLACK}$  ▷ Case 1

$\text{color}[p[x]] \leftarrow \text{RED}$  ▷ Case 1

LEFT-ROTATE( $T, p[x]$ ) ▷ Case 1

$w \leftarrow \text{right}[p[x]]$  ▷ Case 1

**if**  $\text{color}[\text{left}[w]] = \text{BLACK}$  and  $\text{color}[\text{right}[w]] = \text{BLACK}$

**then**  $\text{color}[w] \leftarrow \text{RED}$  ▷ Case 2

$x \leftarrow p[x]$  ▷ Case 2

**else if**  $\text{color}[\text{right}[w]] = \text{BLACK}$

**then**  $\text{color}[\text{left}[w]] \leftarrow \text{BLACK}$  ▷ Case 3

$\text{color}[w] \leftarrow \text{RED}$  ▷ Case 3

RIGHT-ROTATE( $T, w$ ) ▷ Case 3

$w \leftarrow \text{right}[p[x]]$  ▷ Case 3

$\text{color}[w] \leftarrow \text{color}[p[x]]$  ▷ Case 4

$\text{color}[p[x]] \leftarrow \text{BLACK}$  ▷ Case 4

$\text{color}[\text{right}[w]] \leftarrow \text{BLACK}$  ▷ Case 4

LEFT-ROTATE( $T, p[x]$ ) ▷ Case 4

$x \leftarrow \text{root}[T]$  ▷ Case 4

**else** (same as **then** clause with “right” and “left” exchanged)

$\text{color}[x] \leftarrow \text{BLACK}$

# RB deletion fixup

**Idea:** Move the extra black up the tree until

- $x$  points to a red & black node  $\Rightarrow$  turn it into a black node,
- $x$  points to the root  $\Rightarrow$  just remove the extra black, or
- we can do certain rotations and recolorings and finish.

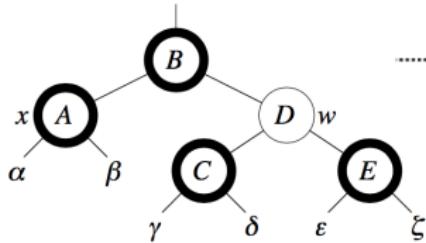
Within the **while** loop:

- $x$  always points to a nonroot doubly black node.
- $w$  is  $x$ 's sibling.
- $w$  cannot be  $nil[T]$ , since that would violate property 5 at  $p[x]$ .

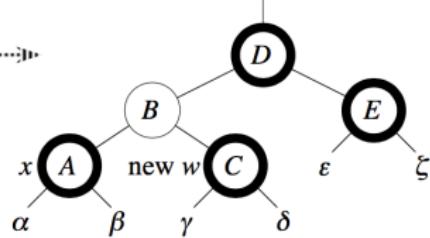
There are 8 cases, 4 of which are symmetric to the other 4. As with insertion, the cases are not mutually exclusive. We'll look at cases in which  $x$  is a left child.

# Case 1

**Case 1:**  $w$  is red



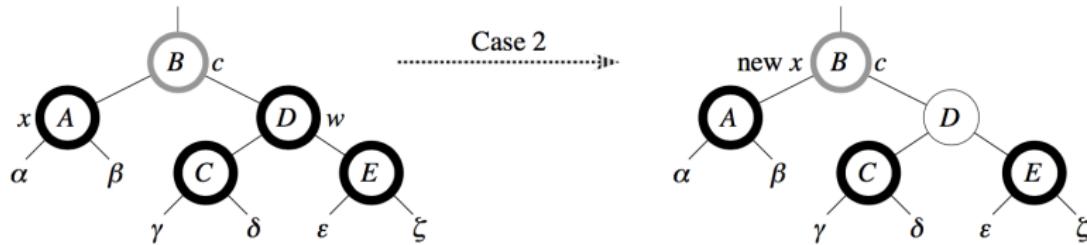
Case 1



- $w$  must have black children.
- Make  $w$  black and  $p[x]$  red.
- Then left rotate on  $p[x]$ .
- New sibling of  $x$  was a child of  $w$  before rotation  $\Rightarrow$  must be black.
- Go immediately to case 2, 3, or 4.

## Case 2

**Case 2:**  $w$  is black and both of  $w$ 's children are black

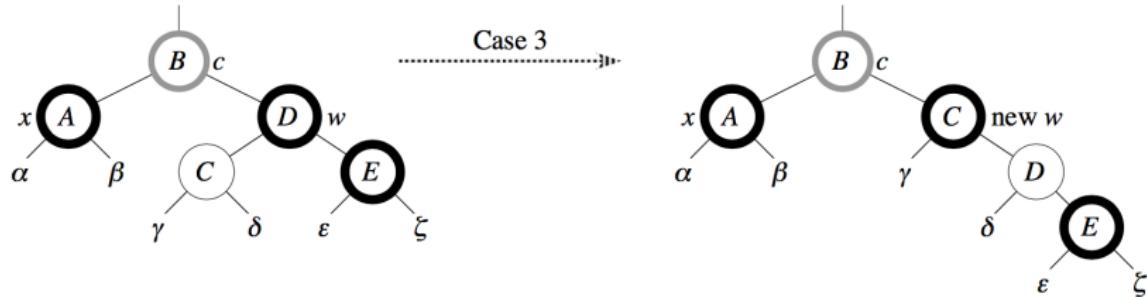


[Node with gray outline is of unknown color, denoted by  $c$ .]

- Take 1 black off  $x$  ( $\Rightarrow$  singly black) and off  $w$  ( $\Rightarrow$  red).
- Move that black to  $p[x]$ .
- Do the next iteration with  $p[x]$  as the new  $x$ .
- If entered this case from case 1, then  $p[x]$  was red  $\Rightarrow$  new  $x$  is red & black  $\Rightarrow$  color attribute of new  $x$  is RED  $\Rightarrow$  loop terminates. Then new  $x$  is made black in the last line.

# Case 3

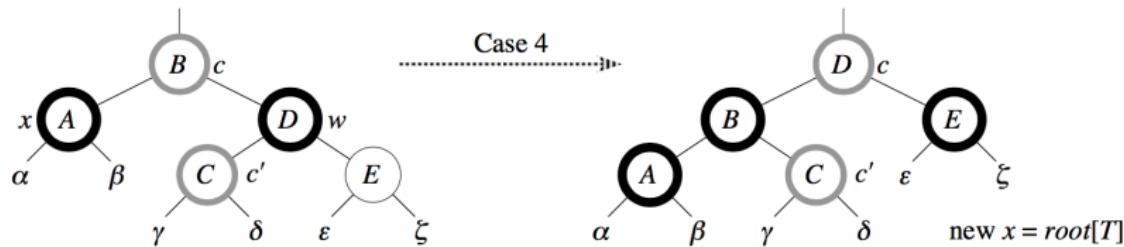
**Case 3:**  $w$  is black,  $w$ 's left child is red, and  $w$ 's right child is black



- Make  $w$  red and  $w$ 's left child black.
- Then right rotate on  $w$ .
- New sibling  $w$  of  $x$  is black with a red right child  $\Rightarrow$  case 4.

## Case 4

**Case 4:**  $w$  is black,  $w$ 's left child is black, and  $w$ 's right child is red



[Now there are two nodes of unknown colors, denoted by  $c$  and  $c'$ .]

- Make  $w$  be  $p[x]$ 's color ( $c$ ).
- Make  $p[x]$  black and  $w$ 's right child black.
- Then left rotate on  $p[x]$ .
- Remove the extra black on  $x$  ( $\Rightarrow x$  is now singly black) without violating any red-black properties.
- All done. Setting  $x$  to root causes the loop to terminate.

# Time analysis

## Analysis

$O(\lg n)$  time to get through RB-DELETE up to the call of RB-DELETE-FIXUP.

Within RB-DELETE-FIXUP:

- Case 2 is the only case in which more iterations occur.
  - $x$  moves up 1 level.
  - Hence,  $O(\lg n)$  iterations.
- Each of cases 1, 3, and 4 has 1 rotation  $\Rightarrow \leq 3$  rotations in all.
- Hence,  $O(\lg n)$  time.